



Department of Informatics
Technical University of Munich



TECHNICAL UNIVERSITY OF MUNICH

DEPARTMENT OF INFORMATICS

MASTER'S THESIS IN INFORMATICS

Efficient and Accurate Hop-by-Hop Capacity Estimation

Bakar Andguladze

TECHNICAL UNIVERSITY OF MUNICH

DEPARTMENT OF INFORMATICS

Master's Thesis in Informatics

**Efficient and Accurate Hop-by-Hop Capacity
Estimation**

**Effiziente und Genaue Hop-by-Hop
Kapazitätsabschätzungen**

Author:	Bakar Andguladze
Supervisor:	Prof. Dr.-Ing. Georg Carle
Advisor:	Simon Bauer, M.Sc. Benedikt Jaeger, M.Sc.
Date:	November 15, 2021

I confirm that this Master's Thesis is my own work and I have documented all sources and material used.

Garching, November 15, 2021

Location, Date

Signature

ABSTRACT

Abstract of the thesis will be written afterwards

CONTENTS

1	Introduction	1
1.1	Motivation	1
1.2	Research Questions	2
1.3	Outline	3
2	Background	5
2.1	Terminology	5
2.2	Network Layers	6
2.3	TCP	6
2.4	UDP	6
2.5	ICMP	6
2.6	Traffic Generation	6
2.7	Capacity Estimation	7
2.7.1	End-to-End vs Hop-by-Hop Capacity Estimation	7
2.7.2	Active vs Passive Capacity Estimation	7
3	Related Work	9
3.1	Existing Approaches	9
3.1.1	Variable Packet Size Probing	9
3.1.2	Packet Pair/Train Probing	9
3.1.3	Pathchar	9
3.1.4	PPrate	9
4	Implementation	11
4.1	Analysis	11
4.2	Approach	11
4.3	Traffic Generation	11
4.4	Capturing the Traffic and Estimating Capacities	12

5	Evaluation	13
5.1	Test Environment and Setup	14
5.2	Evaluation in Empty Network	15
5.2.1	Path Length	17
5.2.2	Packet Size	19
5.2.3	Train Length	21
5.2.4	Optimal Intrusion	23
5.2.5	Capacity Range	23
5.2.6	Packet Loss	25
5.2.7	ICMP Rate Limiting	26
5.2.8	Summary	28
5.3	Evaluation during Cross-Traffic	29
5.3.1	Packet Size	29
5.3.2	Train Length	29
5.3.3	Optimal Intrusion during Cross-Traffic	29
5.3.4	Path Length	29
5.3.5	Capacity Range	29
5.3.6	Packet Loss	29
5.3.7	ICMP Rate Limiting	29
5.3.8	Summary	29
5.4	Data Replication	29
6	Conclusion	31
6.1	Evaluation Results	31
6.2	Answers to the Research Questions	31
6.3	Future Work	32
A	Reproducibility	33
A.1	The Source Code Repository	33
A.2	Setup	34
A.3	Running Tests	35
B	List of acronyms	37
	Bibliography	39

LIST OF FIGURES

4.1	The basic idea of the capacity estimation method	11
5.1	The template network topology	15
5.2	The baseline measurement results	16
5.3	The baseline measurement error rate	17
5.4	Path length: 3 Routers	18
5.5	Path length: 63 Routers	18
5.6	Error rates on different path lengths	18
5.7	Comparison of the measurement results with different packet sizes . . .	20
5.8	Error rates by packet size	20
5.9	The impact of the capacity range on accuracy	24
5.10	Error rates by capacity range	24
5.11	The impact of over 15% packet loss on capacity estimation	25
5.12	Effects of ICMP rate limiting on the accuracy	27
5.13	The impact of long packet trains against ICMP rate limit	28
A.1	The architecture of the framework	34

LIST OF TABLES

5.1	Relative error statistics of the baseline measurement	17
5.2	Relative error stats of paths with different lengths.	19
5.3	Relative error rate statistics for packet size	21
5.4	Relative error statistics for packet train length	22
5.5	Relative error statistics for optimal intrusion	23
5.6	Relative error statistics for the ICMP rate limiting	27

CHAPTER 1

INTRODUCTION

This master thesis presents the capacity estimation method for hop-by-hop measurements. Our goal is to create and test a method of calculating the capacity (i.e. maximum transmission rate) in a given network path and locating the narrow link. We will test the developed tool in regard to various metrics, such as packet size, intrusion rate, cross-traffic and flow-interference and finally, analyze the test results in order to conclude whether our approach is actually capable of delivering accurate results efficiently.

1.1 MOTIVATION

As the Internet is becoming increasingly essential part of our day-to-day lives, it is ever more important for the Internet providers to enhance the quality of networks in order to create a better user experience. One of the means to achieve this goal is to have a better picture of the network they aim to improve. Therefore, we are going to create a tool that measures the capacity of the path between two hosts. The intended field of application is, for instance, enhancing the performance of network, traffic analysis, network monitoring, etc.

There are quite a few capacity estimation methodologies and tools available, that will be discussed in chapters below. However, the current State-of-the-Art methods have some significant flaws and limitations regarding our measurement goals, such as

- High intrusion, which can lead to network overload,
- Dependence on ongoing traffic, which can be unpredictable at times,
- Inability to locate the narrow link on the path.

Therefore, our goal is to develop a new solution that tries to minimize the influence of these limitations in regard to our measurement objectives. We will try to implement the least possible intrusion without compromising the accuracy. Also our tool will be able to find the first narrow link of the path by measuring the capacity of each hop in the network.

This new solution will be tested and evaluated in comparison to the results of existing capacity measurement tools, i.e. PPrate implementation by Patryk Brzoza[x], but in contrast to Brzoza's passive approach our tool will be based on an active measurement methodology. Moreover Brzoza's measurement tool was designed to estimate end-to-end capacity, while this thesis is concerned about measuring the capacity of each hop in the network and finding the narrow link, as hop-by-hop measurements provide a better picture of a network and enable to take a closer look at potential issues.

1.2 RESEARCH QUESTIONS

This thesis is supposed to answer the following research questions:

- **How to measure network capacity hop-by-hop?**

In order to measure the path capacity hop-by-hop, we need to estimate capacities to each router on the path until the destination host is reached.

- **How to optimize the trade-off between accuracy and intrusiveness regarding large-scale measurements?**

Certain level of intrusion into the network will be necessary for the measurements. However there is an important factor to consider: too high intrusion could disrupt the traffic in the network and too low could lead to unreliable results. Therefore an optimal middle ground has to be found: What will be the optimal amount of packets to send to each router to get correct results?

- **How robust is the proposed solution regarding the handling of cross-traffic, flow-interference and sudden path parameter changes?**

Real networks are usually quite complex and different challenges might arise when we are trying to measure the path capacity. We need to find out whether our solution is feasible when it faces cross-traffic, flow-interference or when the path parameters suddenly change.

- **Are we able to locate the capacity bottlenecks of a network?**

We are interested to find the location of the weakest link in the given network. This can be achieved by finding the capacities to each hop. The weakest link will

1.3 OUTLINE

This section introduces the structure of the thesis.

Chapter 2 defines the necessary terminology for understanding this thesis. Namely, capacity, TCP, ICMP, Raw sockets, etc.

Chapter 3 describes the related work - what has been done in regard to capacity estimations and some drawbacks and limitations to the existing approaches

Chapter 4 describes the approach and the tool that we have developed to implement it.

Chapter 5 reviews the test setup and test environment in which the experiments are conducted.

Chapter 6 evaluates our approach based on several parameters. It reviews the different factors that might affect the measurements and to what extent. Namely, how packet length influences the measurement results, what is the optimal rate of intrusion in the network and whether the cross-traffic and the flow-interference cause higher inaccuracy. Based on how the proposed approach handles these challenges we can state whether it is reliable or not.

Finally, chapter 7 concludes our thesis and subsequently discusses the future work - what can be done afterwards to further extend our methodology.

CHAPTER 2

BACKGROUND

The following chapter provides the background information and defines the important terminology that is useful to understand our work.

2.1 TERMINOLOGY

Certain terms in our area of research can be used with different meanings, therefore we first need to state that this thesis will be using the definitions provided by Prasad et al. [1] in their paper "Bandwidth estimation: metrics, measurement, techniques and tools", as it appears to be more widespread and accepted, as those are the definitions also used by other researchers at the chair, therefore, it is more practical to use the common language.

Prasad et al.[1] introduce the following three metrics: capacity, bandwidth and bulk transfer capacity(BTC), capacity being the main focus of our work. Moreover, they distinguish between segments and hops. The former being the link at the data link layer (L2) and the latter - the links at the IP layer (L3).

[THIS IS A COPY] A segment normally corresponds to a physical point-to-point link, a virtual circuit, or to a shared access local area network (e.g., an Ethernet collision domain, or an FDDI ring). In contrast, a hop may consist of a sequence of one or more segments, connected through switches, bridges, or other layer-2 devices. We define an end-to-end path from an IP host (source) to another host (sink) as the sequence of hops that connect to

CHAPTER 2: BACKGROUND

CAPACITY

[ALSO COPY] A layer-2 link, or segment, can normally transfer data at a constant bit rate, which is the transmission rate of the segment. For instance, this rate is 10Mbps on a 10BaseT Ethernet segment, and 1.544Mbps on a T1 segment. The transmission rate of a segment is limited by both the physical bandwidth of the underlying propagation medium as well as its electronic or optical transmitter/receiver hardware. At the IP layer a hop delivers a lower rate than its nominal transmission rate due to the overhead of layer-2 encapsulation and framing. Specifically, suppose that the nominal capacity of a segment is

The transmission time for an IP packet of size bytes is

AVAILABLE BANDWIDTH

Another important metric is the available bandwidth of a link or end-to-end path. The available bandwidth of a link relates to the unused, or “spare”, capacity of the link during a certain time period. So even though the capacity of a link depends on the underlying transmission technology and propagation medium, the available bandwidth of a link additionally depends on the traffic load at that link, and is typically a time-varying metric.

2.2 NETWORK LAYERS

2.3 TCP

<https://datatracker.ietf.org/doc/html/rfc793> header diagrams.

2.4 UDP

2.5 ICMP

<https://datatracker.ietf.org/doc/html/rfc792>

2.6 TRAFFIC GENERATION

RAW SOCKETS

<https://www.binarytides.com/raw-sockets-c-code-linux/>

IPERF

CAPTURING TRAFFIC

<https://www.wireshark.org/docs/man-pages/tshark.html> tcpdump

2.7 CAPACITY ESTIMATION

general explanation

2.7.1 END-TO-END VS HOP-BY-HOP CAPACITY ESTIMATION

2.7.2 ACTIVE VS PASSIVE CAPACITY ESTIMATION

CHAPTER 3

RELATED WORK

3.1 EXISTING APPROACHES

3.1.1 VARIABLE PACKET SIZE PROBING

3.1.2 PACKET PAIR/TRAIN PROBING

3.1.3 PATHCHAR

3.1.4 PPRATE

NOTES

Direct TCP to routers is blocked by routers. (find a source to prove this)

This thesis will find out whether the PPrate algorithm is also useful in hop-by-hop estimations.

Maximum Segment Size tcp

blablabla TCP

CHAPTER 4

IMPLEMENTATION

4.1 ANALYSIS

4.2 APPROACH




FIGURE 4.1: The basic idea of the capacity estimation method

4.3 TRAFFIC GENERATION

```
1  int i = 0;
2  while (1)
3  {
4      //Send the packet
5      if (sendto(s, datagram, iph->tot_len, 0, (struct sockaddr *) &sin, sizeof(sin)) < 0)
6      {
7          perror("sendto failed");
8      }
9      //Data sent successfully
10     else
11     {
12         printf("%d\t", i+1);
13         printf ("Packet Sent Length: %d\n", iph->tot_len);
```

CHAPTER 4: IMPLEMENTATION

```
14     }
15
16     // ttl should be incremented after every n packets
17     i++;
18     if(i == packets) // number of packets necessary for measuring
19     {
20         printf("hop_#%d_done\n\n", iph->ttl);
21         iph->ttl++;
22         i = 0;
23     }
24
25     if(iph->ttl == routers+2) // total number of routers + 2
26     {
27         break;
28     }
29 }
```

4.4 CAPTURING THE TRAFFIC AND ESTIMATING CAPACITIES

CHAPTER 5

EVALUATION

This chapter presents the results of the evaluation of our approach. The experiments have been conducted based on several test parameters and their combinations and we will be analyzing the baseline parameter sets as well as the key combinations.

The capacity estimation tools can face different challenges in real networks. We will try to imitate some of these challenges in our custom-built emulated network and find out what kind of impact can they have on measurement accuracy.

Two of the biggest obstacles our framework has to face are the cross-traffic and the ICMP rate limiting, which we will discuss shortly in respective sections.

Overall, we will evaluate the tool based on the path length, the capacity range, the packet size, the packet train length. For a deeper insight we will also apply the ICMP rate limiting and also check the influence of the packet loss on accuracy. The tests are conducted on an empty network as well as with the significant amount of cross-traffic load.

The following are the parameters which we use for the evaluation of the tool. They are manually configured by us for each test suite in order to get a better picture about the quality of our approach and shows us whether it is applicable in real life:

- **Path Length** - the number of routers between the source host and the destination host. (i.e. number of hops minus one)
The default value: 8 routers.
- **Capacity Range** - The range of numbers from which a capacity for each link will be generated
The default value: [10, 100).

- **Packet Size** - The size of each packet that will be sent from the source host to the destination. It represents the total size of the packet including TCP and IP headers(each weighing 20 bytes).
The default value: 1500 bytes (i.e. size of MSS + TCP header + IP header = 1460 + 20 + 20)
- **Packet Train Length** - The amount of packets that will be targeted at each hop
The default value: 300 packets.
- **Packet Loss** - The emulated packet loss on each router.
The default value: 0.
- **ICMP Rate Limit** - The artificial ICMP rate limit that will be assigned to each router.
The default value: 0.
- **Cross Traffic** - The coefficient of the cross traffic load. The Optimal value is recommended to be assigned from 0 to 1.0, the value of 0 meaning an empty network without any cross traffic.

To ensure the reliability on the the test results, we conduct test runs with each test configuration 20 times.

5.1 TEST ENVIRONMENT AND SETUP

For the testing purposes we decided to build a virtual network using the network emulator - Mininet[2]. It provides the necessary tools to create an artificial network and has a very handy Python API that enables us to build custom topologies.

The figure 5.1 represents the template of the network that we have used in our experiments. It consists of the source and the destination hosts, the top and the bottom hosts and the routers that connect them with each other.

As a reader can see from the Figure 5.1, the main hosts (sender and receiver) are connected with a sequence of n routers, each consisting of four interfaces and each having its own subnetwork. The number of routers and, therefore, the number of top and bottom hosts is dynamic and configurable. The routes from and to all hosts are statically configured.

The measurements are conducted based on the main path, which is the one that consists of the routers (1... n) and connects the sender to the receiver. The top and the bottom hosts are used for generating cross-traffic in later experiments. The cross-traffic




FIGURE 5.1: The template network topology

interferes with the main flow of packets and has a significant influence on the accuracy of capacity estimations.

The measuring node is the "Sender", as using our tool only makes sense if the measuring node is the one that generates the traffic. All the necessary traffic is captured at this node with the `tcpdump` command and subsequently analyzed by our program that delivers the final result.

Finally, it is also important to mention that Mininet has certain limitations, namely, it is limited with the computational power of the machine it runs on. Therefore, we have decided to execute the estimation experiments on the testbed of the chair.

Disclaimer: As the routers are all configured the same way in terms of packet loss and ICMP rate limiting, the results from all hops are analyzed together. The tests have shown no difference between the accuracy rates of different indexes of hops, e.g. the accuracy rate of the capacity estimations from the source to R_1 and to R_n is practically the same in equal circumstances.

5.2 EVALUATION IN EMPTY NETWORK

We decided to divide our evaluation into two parts: the section 5.2 describes the measurement results on an empty network and the section 5.3 analyzes the effect of the cross

traffic based on the experiment outcomes on an overloaded network where the flow of our generated traffic is interfered by the cross-traffic.

To have a broad picture of the initial results, we first conducted baseline measurements with the default configuration of the parameters, namely, the path length of 8 routers and the packet size of 1500 bytes in a capacity range of $[10, 100)$, which resulted in accurate estimations.

The graph in the figure 5.2 represents the comparison between the real and estimated capacities and as one can see, the accuracy is high.

Moreover, the figure 5.3 and the table 5.1 represent the error rate of the baseline measurement and statistical details about error rates respectively.




FIGURE 5.2: The baseline measurement results

5.2 EVALUATION IN EMPTY NETWORK




FIGURE 5.3: The baseline measurement error rate

Average	0.67%
Standard Deviation	0.27%
Min Error	0.07%
Max Error	1.55%

TABLE 5.1: Relative error statistics of the baseline measurement

The baseline evaluation shows the the good potential of our approach, however it was conducted with the simplest and likely some of the most promising configuration possible, which is not likely to be the case in real networks.

In the next sections we will dissect each parameter through subsequent experiments and check their impact on the accuracy in order to evaluate how robust and useful our capacity estimation methodology can be in real life.

5.2.1 PATH LENGTH

The first parameter to be analyzed is the path length from the source host to the sink, i.e. number of routers on the path. In order to analyze the impact that a path length can have on the measurement accuracy, we have to execute the test runs on different values of the given parameter.

As the default TTL value in Linux for TCP and ICMP is 64[3], we decided to conduct measurements on arbitrary numbers of routers up to 63, namely 1, 3, 8, 20, 32 and

63. These numbers, although being chosen randomly apart from 1 and 63, paint a good picture of how our estimation tool reacts on networks with different sizes. Based on the results we can assume that the path length does not have an impact on the accuracy of our capacity estimation framework. Figures 5.4 and 5.5 depict the differences between actual and estimated capacity values on networks with path length of 3 and 63 respectively. (We chose to present only these two graphs, as there are practically no relevant differences between the results of different path lengths)

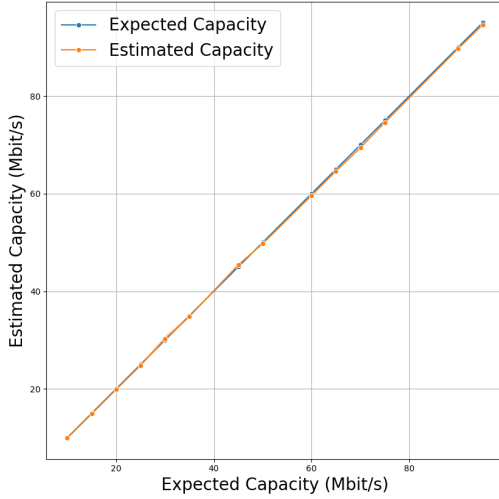


FIGURE 5.4: Path length: 3 Routers

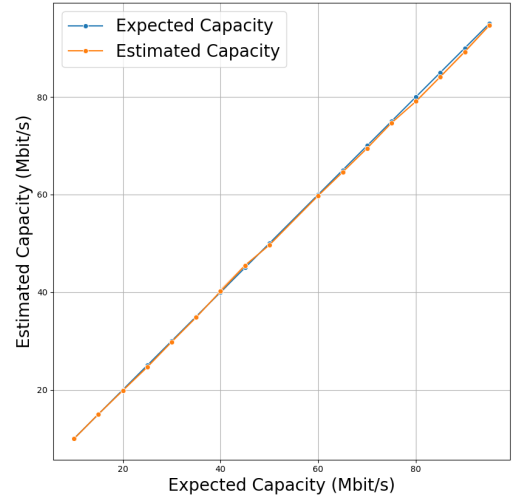


FIGURE 5.5: Path length: 63 Routers




FIGURE 5.6: Error rates on different path lengths

5.2 EVALUATION IN EMPTY NETWORK

The table 5.2 depicts the comparison of relative error rates in more details and as one can observe, the differences are practically irrelevant.

TABLE 5.2: Relative error stats of paths with different lengths.

Path Length	Average	Standard Deviation	Min Error	Max Error
1	0.65%	0.21%	0.23%	1%
3	0.72%	0.5%	0.11%	4.3%
8	0.74%	0.47%	0.01%	2.83%
20	0.75%	0.24%	0.01%	2.87%
32	0.83%	0.48%	0.0001%	4.09%
63	1.23%	2.53%	0.01%	42.57%

CONCLUSION

Judging by the results of the experiments focusing on the path length conducted in an empty network we can conclude that the number of hops does not affect the accuracy of the estimation tool. At this point it can be safely assumed that when the cross-traffic is out of the picture, our framework can estimate capacity accurately no matter the amount of hosts on the path. Although there is still a small likelihood of major inaccuracies in measurements, such as $\approx 43\%$ relative error, it does not affect the overall picture, as it can be seen in the graphs.

After testing the path length we will further use 8 routers for the subsequent estimations as a default value.

5.2.2 PACKET SIZE

Our next target parameter is the size of each packet in our generated traffic. As the packet size along with the amount of packets that have to be injected into the network can be the cause of the network overload, we have to find an optimal packet size that most importantly does not cost us the accuracy and also the load imposed on a target network remains minimal. In other words, the goal is to find the least packet size value that will deliver the accurate results.

We launched the tests with the packets of Maximum Transmission Unit (MTU) size, i.e. 1500 bytes and with the packets as small as 100 bytes. As depicted before, the former delivers highly accurate results, the latter, however, has shown an average relative error rate of $\approx 30\%$, thus we decided to narrow down between the two values and find out at which packet size does relative error rate remain consistently low, i.e. what is the minimal packet size that matches the accuracy of 1500 bytes.

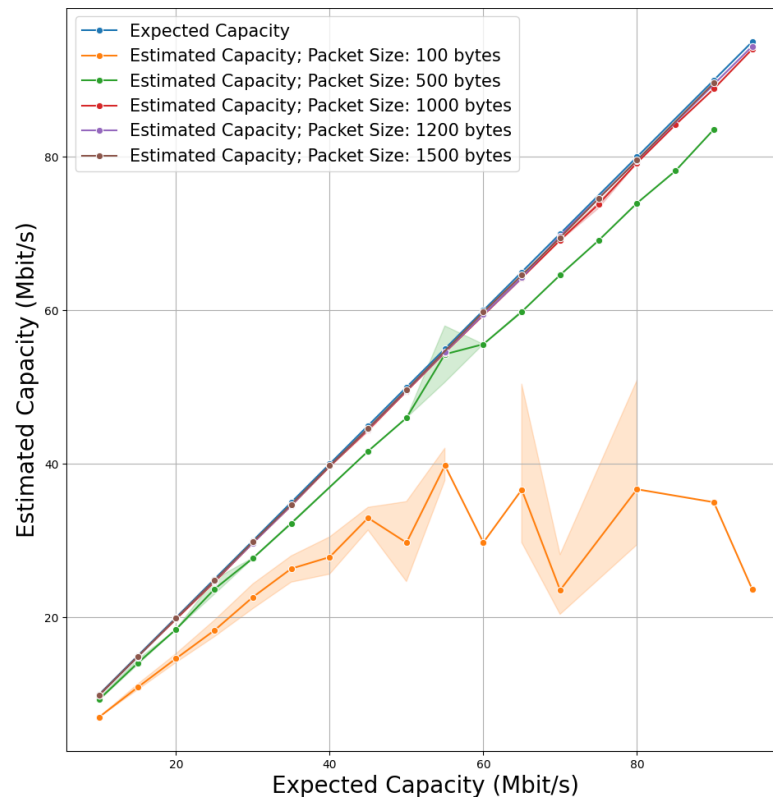


FIGURE 5.7: Comparison of the measurement results with different packet sizes




FIGURE 5.8: Error rates by packet size

5.2 EVALUATION IN EMPTY NETWORK

Subsequently we tested packet sizes of 500, 1000 and 1200 bytes which showed drastic improvements in estimation accuracy. The results of these experiments are depicted in the figure 5.7 the respective error rates in the figure 5.8.

Depending on a researcher's threshold of the tolerable relative error rate the minimal packet size can vary. Thus we launched several more test runs with packet sizes between 500 and 1000 bytes in order to narrow down even more and to provide more detailed research.

The table 5.3 shows the results of all the tests on packet size.

TABLE 5.3: Relative error rate statistics for packet size

Packet Size	Average	Standard Deviation	Min Error	Max Error
100	30.23%	10.97%	6.29%	75.24%
500	7.96%	3.11%	1.12%	41.57%
600	2.11%	0.66%	0.1%	4.85%
700	1.8%	0.64%	0.01%	5.92%
800	1.53%	0.47%	0.14%	4.6%
900	1.3%	0.54%	0.07%	5.42%
1000	1.15%	0.48%	0.0%	3.57%
1200	0.93%	0.25%	0.1%	1.72%
1500	0.77%	0.32%	0.13%	2.98%

CONCLUSION

The estimations show that the impact of the packet size on the accuracy is huge and the packet size of 1500 bytes delivers the most accurate results. The accuracy declines along with the decreasing packet size. However, if we take, for example, less than 2% as an average relative error rate that can be tolerated, we can decrease packet size down to 750 bytes (the half of the original 1500 bytes) and still get highly accurate results. We should expect, though, that this number might not be sufficient in networks where the cross traffic is the case. This will be discussed in another section.

5.2.3 TRAIN LENGTH

In the original paper of PPrate[4] En-Najjary and Urvoy-Keller state that the second version of the PPrate algorithm has good estimation results with as many as 300 IAT samples. Our previous estimations have also shown that the train length of 300 packets (299 IATs if packet loss is 0) can potentially be enough for accurate results with less than 1% error rate.

But in this thesis we are aiming for the minimal intrusion. Therefore we are interested in decreasing the packet train lengths without compromising accuracy and finding out

the least number of packets that will be provide the estimations as accurate as we have achieved with 300 packets.

The version of the Python implementation of PPrate by Brzoza[5] that we are using for our estimations does not have a constraint on the amount of IAT samples. Thus we experimented on as small numbers of packets as possible.

We conducted tests with different train lengths and the outcome was unexpected: The smallest train length that gave us accurate estimations was only 7 packets, meaning only 6 IATs were sufficient for PPrate to estimate capacity accurately. The algorithm was unable to calculate capacities with less than 6 IATs.

Nevertheless there is one more factor to consider. In our experiments conducted in Mininet we could observe that the packet trains with less than 50 packets had almost no packet loss and increasing the train length often resulted in higher rate of loss without explicitly configuring packet loss. E.g. after 20 experiments with the trains of 7 packets the loss was 0.

TABLE 5.4: Relative error statistics for packet train length

Train Length	Average	Standard Deviation	Min Error	Max Error
7	1.42%	1.61%	0.01%	10.09%
10	1.25%	1.3%	0.03%	10.78%
20	1.1%	1.0%	0.05%	5.84%
50	0.81%	0.47%	0.01%	4.12%
100	0.78%	0.41%	0.03%	2.95%
250	0.74%	0.3%	0.07%	2.76%
500	0.53%	0.38%	0.05%	2.94%
1000	0.49%	0.43%	0.04%	3.39%
5000	0.42%	0.36%	0.02%	2.72%

The table 5.4 shows shows the statistics of relative error rate depending on the train length. It suggests that with the increase of the train length average error rate decreases, but the difference in estimation results between 7 packet train and 5000 packet train is only 1%.

CONCLUSION

The test results are showing that PPrate is able to provide high accuracy with a very small number of packets in a packet train when we are running tests on an empty network. Although more importantly it is notable that longer packet trains provide more accurate results.

These experiments, however, are not enough to assess that we can get the similar results

5.2 EVALUATION IN EMPTY NETWORK

in real networks where the cross-traffic is also an issue. This factor will be tested separately in section 5.3.

5.2.4 OPTIMAL INTRUSION

Based on our experiments on empty network we came to a conclusion that the packet size is more crucial parameter to the capacity estimation rather than packet train length. It is necessary that the packet size equals at least 700 bytes in order to keep the relative error rate reasonable, for example, below 2% on average.

When it comes to the train length, however, the situation is more flexible here. Meaning, PPrate can deliver accurate results in empty network with as few as 10 packets. We combined these values and conducted the tests with 10 packet trains and 700 bytes packet size. It resulted in 3.36% average relative error rate, therefore we increased both parameters several more times until the average error dropped below 2%.

The table 5.5 shows the results of those experiments.

TABLE 5.5: Relative error statistics for optimal intrusion

Packet Size	Train Length	Average	Standard Deviation	Min Error	Max Error
700	10	3.36%	5.33%	0.04%	48.69%
700	50	2.19%	1.43%	0.08%	10.93%
800	10	6.7%	6.4%	0.01%	32.91%
800	30	1.81%	1.03%	0.11%	6.72%
1000	30	1.32%	0.44%	0.06%	3.75%

Assuming that real networks are not usually empty, we can hypothesize that in practice higher intrusion will be necessary for relatively precise estimations. We will further research this topic in section 5.3

5.2.5 CAPACITY RANGE

One of the most important parameters is the capacity range of the links that we conduct our measurements on. So far we have seen that our tool can handle the links with capacities between 10 and 100 Mbits. The real networks can contain hops with higher capacities though, therefore it is necessary to research that aspect as well.

We tested the tool on several different networks with varying capacity ranges, such as [10, 100), as well as [100, 400), [400, 600) and [600, 1000) and as it appears, our methodology has a significant limitation when measuring the networks that contain links with relatively high capacities. Figures 5.9 and 5.10 show that the inaccuracy starts increasing with higher capacities.




FIGURE 5.9: The impact of the capacity range on accuracy

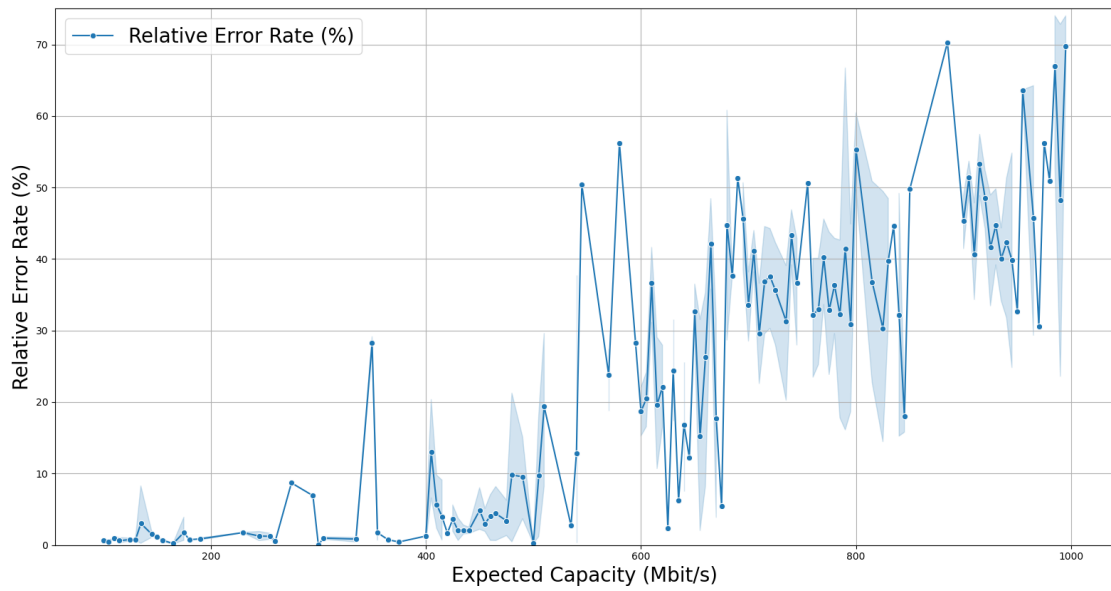


FIGURE 5.10: Error rates by capacity range

CONCLUSION

In conclusion, we can deduce that the accuracy of the tool greatly depends on the capacity range. We can argue that on an empty network the accuracy is good for the paths that have the capacity value up to ≈ 250 Mbits. However the pattern emerges that capacities above 250 Mbits are not underestimated below 250. We can assume that the estimation results below 250 are accurate. We will further observe this pattern with cross traffic in section 5.3

5.2.6 PACKET LOSS

It is a very commonly occurred event in everyday internet traffic when data packets fail to reach their destination. This event is referred as packet loss and there can be different reasons that cause it. According to A. S. Gillis[6] it can be caused by overloaded network, software and/or hardware problems, etc.




FIGURE 5.11: The impact of over 15% packet loss on capacity estimation

In order to make sure that our approach can cope with it and deliver the accurate estimations despite the lack of lost packets we tested our tool against this issue as well. We chose the data samples where the total packet loss in each test run was above 15%.

These samples had on average 35% total packet loss. Despite such a high loss the Figure 5.11 shows that the estimation accuracy was not compromised.

According to the test results so far, we can assume that the random packet loss does not affect the accuracy to extent that affects accuracy, because most of the time it happens to the larger packet trains and the amounts of the ICMP packets that reach the source host is high enough for PPrate to estimate capacities precisely enough.

Although there is a special case of packet loss that represents a serious challenge to our estimation methodology, that is ICMP rate limiting. In contrast with the end-to-end estimation, during which the measurements are based on inter-arrival-times of TCP ack signals, ICMP messages are prone to a huge amount of packet loss due to the ICMP rate limits of internet routers.

We will further discuss this issue in the next section.

5.2.7 ICMP RATE LIMITING

One of the biggest obstacles active ICMP-based measurement tools can face is ICMP rate limiting as it causes a huge rate of packet loss and also the packets that reach destination might not be suitable for capacity measurements. According to the Linux man page[7] it works in the following way: there are two parameters that have to be configured for ICMP rate control: `icmp_ratemask` and `icmp_ratelimit`. The first parameter consists of 19 bits each determining which types of ICMP should be affected by the rate limit (each bit corresponds to a specific ICMP type) and the second one sets the limit of the ICMP responses. `icmp_ratelimit` indicates the minimum time interval between ICMP responses and it is measured in milliseconds. The value range is $[1, 1000]$ [8] and in Linux it has a default value of 1000.

All of our previous measurements were conducted with disabled `icmp_ratemask`. We have, however, also tested the effects of the rate limiting on our tool and tried to find out whether this can be a defining factor in the applicability of our methodology.

As the Figure 5.12 and the table 5.6 show the results are very unpredictable and thus unreliable.

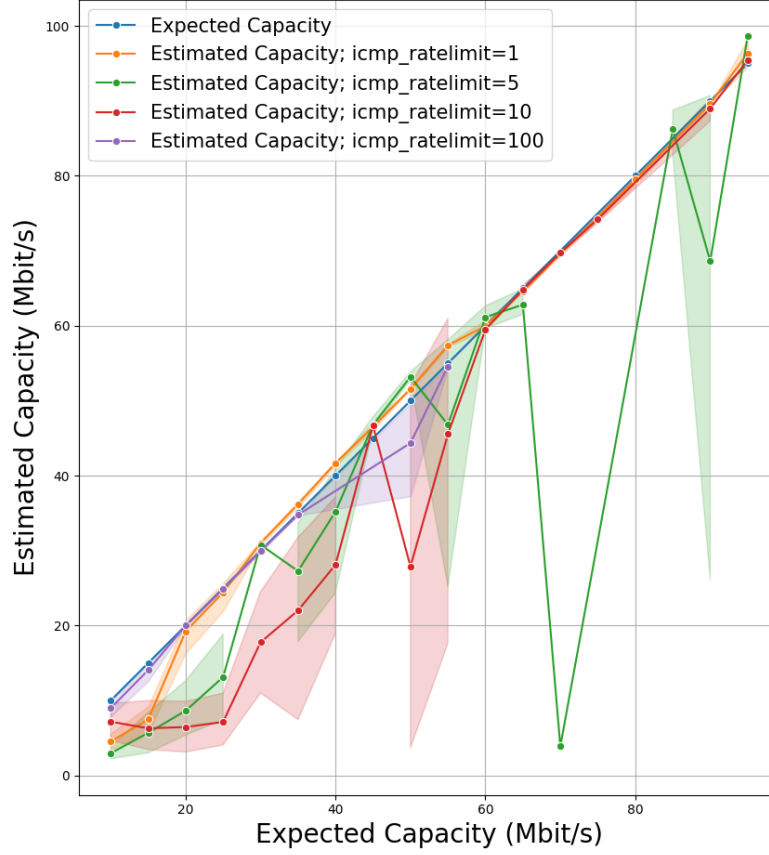


FIGURE 5.12: Effects of ICMP rate limiting on the accuracy

TABLE 5.6: Relative error statistics for the ICMP rate limiting

Train Length	Average	Standard Deviation	Min Error	Max Error
1	22.19%	31.59%	0.02%	84.94%
5	50.45%	41.24%	0.1%	94.45%
10	48.24%	44.35%	0.13%	96.09%
100	5.79%	19.79%	0.06%	96.28%

Nevertheless we have made one observation, that resulted in some improvements in terms of accuracy: the increase in train length provided more accurate results, as despite the huge packet loss rate, the number of 'surviving' packets was high enough to drastically improve the accuracy rate.

The Figure 5.13 presents the results of the estimation with the train length of 5000 packets and `icmp_ratelimit` of 1000 milliseconds.




FIGURE 5.13: The impact of long packet trains against ICMP rate limit

Further research is required to find better solutions to this important problem, which it is out of scope of this thesis, however there are some efforts made in this direction, namely the papers by Ravaioli et al. [9] and H. Guo and J. Heidemann[10]

5.2.8 SUMMARY

To summarize this section, we have examined our capacity estimation technology in empty networks and tested it against different parameters. The results so far have shown that the tool provides high accuracy on any amount of hops with small packet trains and packet size of about 800 bytes. The tool has also shown some downsides, namely the inaccuracy is significant on high capacity networks and it is not reliable when the routers on the path have limits on ICMP response rates.

In the next section we will try to test our framework in a close to real life scenario, i.e. apply artificial cross traffic to each router on the path and observe its effects on accuracy.

5.3 EVALUATION DURING CROSS-TRAFFIC

We have conducted the previous experiments on empty networks without any cross-traffic. However this is not the case in the real Internet that we use on daily basis. Cross traffic seems to affect the test results significantly. What does 1.0 mean? How our cross traffic works. - Formula of cross traffic

5.3.1 PACKET SIZE

5.3.2 TRAIN LENGTH

5.3.3 OPTIMAL INTRUSION DURING CROSS-TRAFFIC

5.3.4 PATH LENGTH

5.3.5 CAPACITY RANGE

5.3.6 PACKET LOSS

5.3.7 ICMP RATE LIMITING

5.3.8 SUMMARY

5.4 DATA REPLICATION

One option to try to see if it can help

CHAPTER 6

CONCLUSION

In this thesis we have implemented a hop-by-hop capacity estimation methodology and evaluated it based on various parameters.

The following sections will shortly summarize our work throughout the timespan of the thesis and address all the concerns and questions that we had in the beginning.

6.1 EVALUATION RESULTS

Sum up whether this method can actually be used in practice

Like other active capacity estimation tools, ours also depends on the network that has to be measured. Less loaded the network is, the better are the results.

6.2 ANSWERS TO THE RESEARCH QUESTIONS

After the careful and thorough research we can already answer the questions we were aiming to address in the beginning of this thesis:

- **How to measure network capacity hop-by-hop?**

We have developed a way to estimate capacity hop-by-hop by sending ttl-exceeding packets to all the routers on the path to the destination.

- **How to optimize the trade-off between accuracy and intrusiveness regarding large-scale measurements?**

- **How robust is the proposed solution regarding the handling of cross-traffic, flow-interference and sudden path parameter changes?**
- **Are we able to locate the capacity bottlenecks of a network?**

The answer to this question greatly depends on the accuracy of the hop-by-hop estimation. As we have seen in the Evaluation chapter, there are situations when we get accurate results and there are cases when the results are somewhat inaccurate or not reliable at all. In case of the reliable results, the answer will be 'yes'. We can locate the capacity bottlenecks based on the first smallest capacity value from the sequence of results.

There is a catch, however: if the path consists of multiple hops and there are several bottlenecks, we are able to locate only one, namely the first one. For example, if there are 15 hops on the path and the bottleneck is located on the first hop, but the capacities of several other links equals to that of the first link, we will not be able to see it. In such situations, we can only assume that the subsequent capacities are more than or equal of the capacity of the first hop.

6.3 FUTURE WORK

As our evaluation has shown, there still is a lot of room for improvement for our methodology.

First and foremost, we have not conducted any test runs on a real network, as the performance of Mininet greatly depends on the machine it runs on and this could compromise the accuracy of the tool. Therefore, internet measurements should be performed in order to see how our tool handles the real traffic in the Internet.

Moreover, further research is required to handle the inaccuracy problem with cross-traffic as it is the key factor when it comes to the real-life usage of our framework.

Fader[10]

CHAPTER A

REPRODUCIBILITY

This chapter describes the source code repository and provides the information and instructions of how to use, alter and/or improve the current version of the Hop-by-hop measurement framework.

A.1 THE SOURCE CODE REPOSITORY

The source code repository for this thesis is located on the GitLab instance belonging to the TUM. The `ma-andguladze/App` directory contains the whole source code.

The tool consists of the following files:

- `PPrate.py` - The pprate algorithm implementation by Patryk Brzoza[5]
- `TrafficGenerator.c` - As the name suggests, this program generates TCP traffic from one host to another. It takes the IP addresses of two hosts as arguments and creates the traffic via raw sockets. The file should be compiled before the first run.
- `mininet_topo.py` - Builds an instance of the network in Mininet considering all the necessary parameters passed, conducts the traffic generation, captures the traffic and saves the results in pcap files.
- `prepare_test.py` - Contains several helper functions, most notably the config file parser and the capacity generator.
- `process_tcp_csv.py` - Converts the Calculates the end-to-end capacity of the network. Implementation by Brzoza[5]

- `process_icmp_csv.py` - Converts the `icmp.pcap` file into csv and delivers the hop-by-hop capacity estimations based on the latter
- `run_test.py` - The main file that bootstraps the tool and runs the experiments in one setting. It requires a JSON configuration file to be passed as an argument with the help of which it is possible to manipulate several key parameters necessary for testing.




FIGURE A.1: The architecture of the framework

A.2 SETUP

In order to run the test measurements with our tool, several dependencies and software should be installed on the Linux machine.

The experiments conducted by us were run only on Mininet - a network emulator tool which creates a realistic virtual network of virtual hosts, switches, controllers and links and runs a real kernel and application code[2]. To run the experiments on Mininet, it is necessary to install the emulator first. It is highly recommended to install Mininet from the source GitHub repository instead of packages, as the Mininet homepage[11] states that packages could give an older version.

Mininet installation:

```

1  $ git clone git://github.com/mininet/mininet
2  $ mininet/util/install.sh [options]

```

A.3 RUNNING TESTS

The default branch normally will install the latest stable version, however in case a user wishes to have any other version, it is also possible to check out the branch with the following script before running `install.sh`:

```
1 $ cd mininet
2 $ git tag # list available versions
3 $ git checkout -b mininet-2.3.0 2.3.0 # or any desired version from the list
4 $ cd ..
```

Note that Mininet must run as a root.

As for the rest, the following dependencies are required to be installed to run experiments:

- pandas
- NumPy
- SciPy
- tshark
- iPerf

For setting up these and compiling the `TrafficGenerator.c` program, run `install.sh` script in the `ma-andguladze/App` directory.

A.3 RUNNING TESTS

The measurements must be run with a JSON config file containing all the parameters that are interesting and important for our measurements. Manipulating these parameters gives us a better picture about strengths and weaknesses of our approach.

The following script executes the experiment:

```
1 $ sudo python run_test.py <config file>
```

The following JSON serves as an example of a configuration file used for experiments. It contains default parameter values, each of which will be manipulated respectively in upcoming sections:

```
1 {
2   "topo_size": 3, // path length
3   "capacity_range": [10, 100],
4   "capacity_delta": 5
5   "packet_size": 1400,
6   "packets_per_hop": 300,
7   "icmp_ratelimit": 0,
8   "packet_loss": 0,
9   "cross_traffic": 0.0,
```

```
10     "output": "results/results.csv"  
11 }
```

***Disclaimer:** `packet_loss` is a probability, therefore it does not always cause the packet loss during test runs; On the other hand, regular test runs with `packet_loss = 0` can result in lost packets.*

CHAPTER B

LIST OF ACRONYMS

MSS	Maximum Segment Size.
MTU	Maximum Transmission Unit.

BIBLIOGRAPHY

- [1] R. Prasad, C. Murray, C. Dovrolis, and K. Claffy, “CSMA/CA Bandwidth Estimation: Metrics, Measurement Techniques, and Tools”, Dec. 2003.
- [2] *Mininet Home*, <http://mininet.org/>, [Online; accessed 2021].
- [3] S. Iveson, *IP Time to Live (TTL) and Hop Limit Basics*, <https://packetpushers.net/ip-time-to-live-and-hop-limit-basics/>, May 2019.
- [4] T. En-Najjary and G. Urvoy-Keller, “PPrate: A Passive Capacity Estimation Tool”, May 2006.
- [5] P. Brzoza, “Implementation and Evaluation of Passive Capacity Estimation”, B.Sc. Thesis, Technische Universität München, Munich, Germany, 2019.
- [6] Alexander S. Gillis, *Packet Loss*, <https://www.techtarget.com/searchnetworking/definition/packet-loss/>, [Online; accessed 11.2021].
- [7] *icmp - Linux IPv4 ICMP kernel module*, <https://man7.org/linux/man-pages/man7/icmp.7.html>, [Online; accessed 11.2021].
- [8] Juniper Networks, *icmp - Error Message Rate Limit*, <https://www.juniper.net/documentation/us/en/software/junos/transport-ip/topics/ref/statement/icmp-edit-chassis.html/>, [Online; accessed 11.2021].
- [9] R. Ravaioli, G. Urvoy-Keller, and B. Chadi, “Characterizing ICMP Rate Limitation on Routers”, 2015.
- [10] H. Guo and J. Heidemann, “Detecting ICMP Rate Limiting in the Internet(Extended)”, Apr. 2017.
- [11] *Mininet Installation*, <http://mininet.org/download/>, [Online; accessed 11.2021].